

Hello, would you like to hear a TCP joke?

Yes, I'd like to hear a TCP joke.

OK, I'll tell you a TCP joke.

OK, I'll hear a TCP joke.

Are you ready to hear a TCP joke?

Yes, I am ready to hear a TCP joke.

OK, I'm about to send the TCP joke. It will last 10 seconds, it has two characters, it does not have a setting, it ends with a punchline.

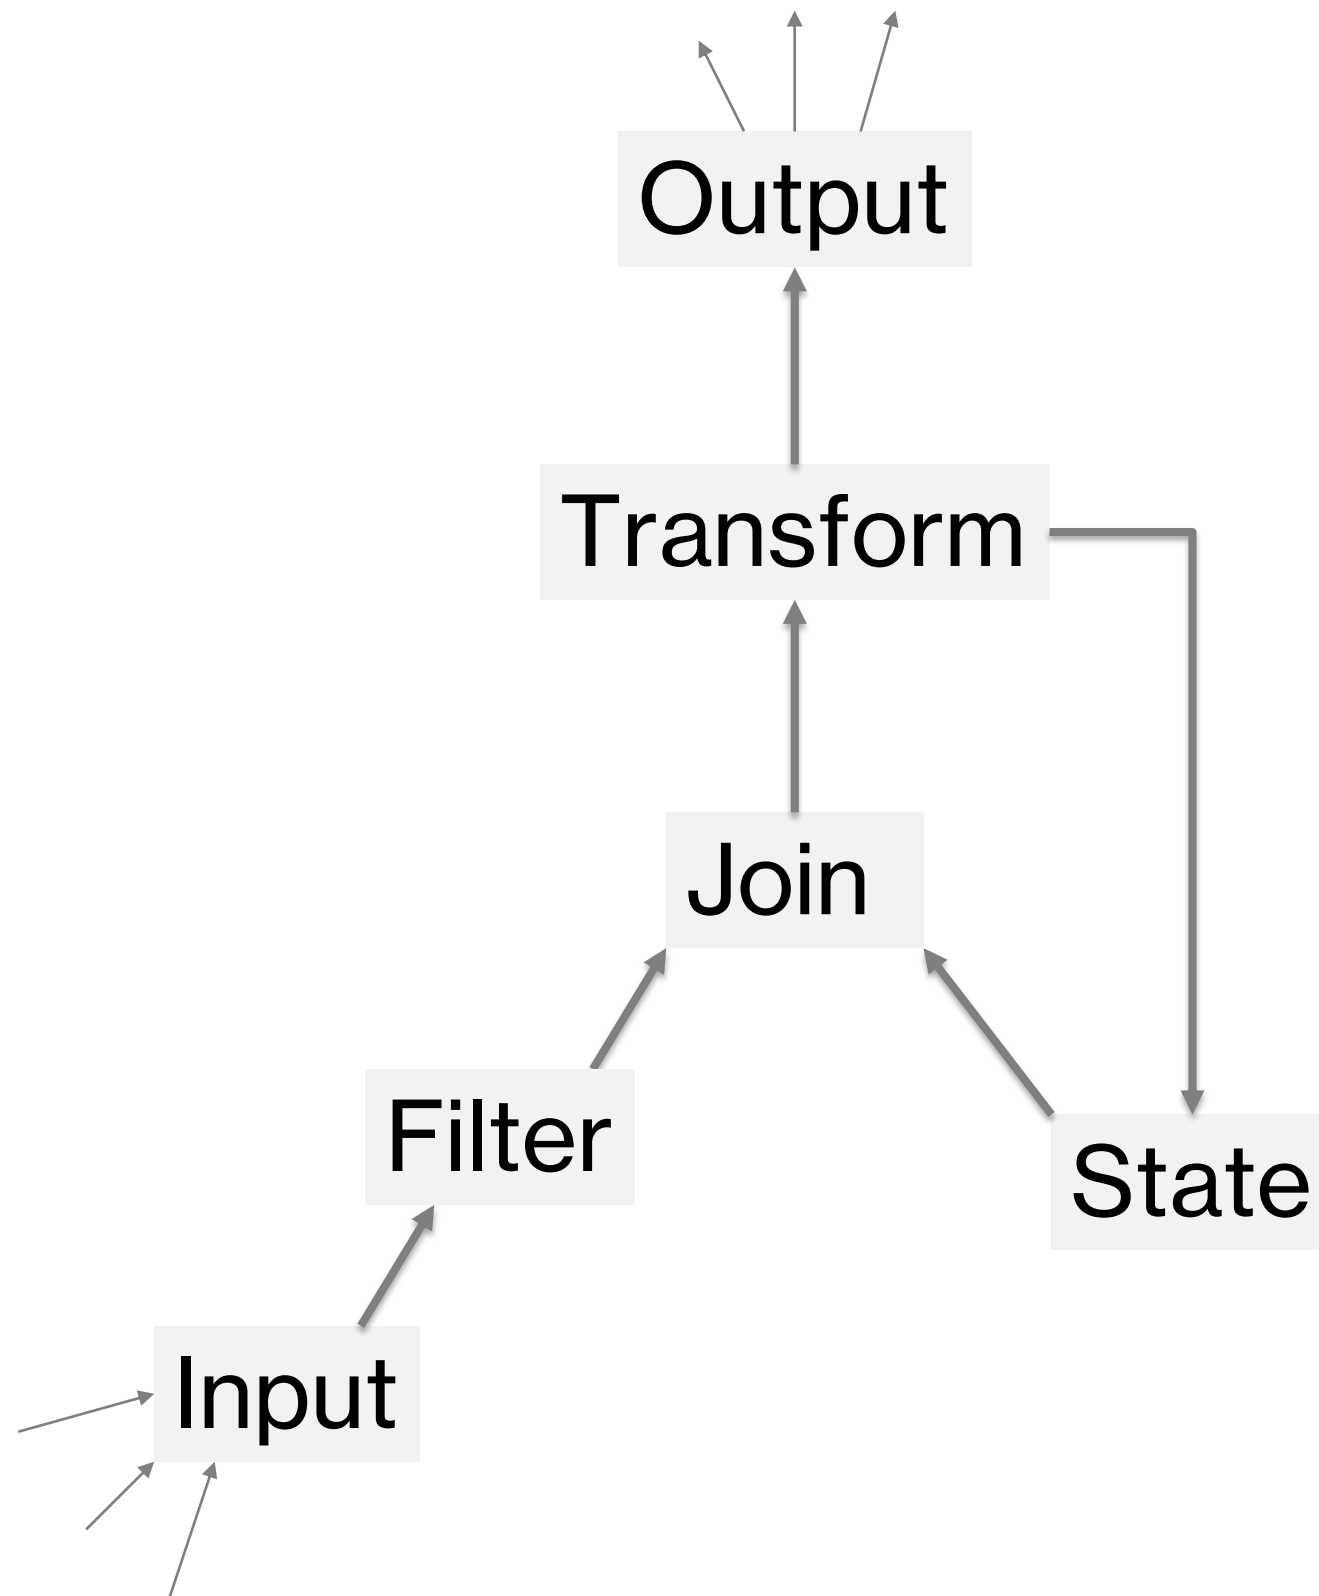
OK, I'm ready to hear the TCP joke that will last 10 seconds, has two characters, does not have a setting and will end with a punchline.

I'm sorry, your connection has timed out... ..Hello, would you like to hear a TCP joke?

# Distributed systems

# 1: Processes

# The generalized process



`output = fn1(input, state)`

`state = fn2(input, state)`

# Observations I: internal

Two flows: Control and data

For control: Polling is "downward", events are "upward"

State feedback loop captures time

Discrete time, in the case of events

This is a generalized "state machine"

Each process has exclusive ownership of its data

The process accepts or rejects transformations

Transformation functions can be pure

(Though expressing "list of actions" can be hard)

# Observations II: connections

The output of one process is the input of another

This can be local or distributed

reliable or unreliable messages

Messages are values (copies of data)

this is what lets the process reject transformations

Can limit access to inputs with channel-like mechanisms

not "all or nothing" like OO-objects or actors

Sending/receiving is destructive

# Observations III: conclusions

1: The process must be stateful to be useful

2: Messages must carry transformations of state to be useful

Statefulness is usefulness

1+2: Losing/reordering messages is losing/scrambling state

There is no general abstraction for distribution

## 2: Hierarchies of time



# Flavors of unreliability in messages

## Loss

Time doesn't happen

## Duplicates

Time happens twice

## Reordering

Time happens in the wrong sequence

Common: Progression of time is not constant

# Using the message-process duality

## Solutions to non-constant time:

- Discard non-constant rate messages (lockstep)

  - This is the impossible general abstraction

- Discard non-monotonic messages (clocks or counters)

  - Often viable: eg only need the "latest update"

- Accept a select few non-monotonic messages

  - Possible: Carefully controlled reverse-time (eg "resets")

- Accept all messages (and all inconsistencies)

  - Usually not viable: wrong-order inputs can create wrong-order outputs

# The tradeoff in the solutions

Consistency at the cost of protocol

(also at the cost of availability and partition tolerance)

This protocol is code that must run in the process

Responsibility possibly shared among both sender and receiver

Protocol is limitation

Limits the allowable process outputs and state transformations

(The "state machine" gets more states)

# 3: Protocols

# What consistency means

In distributed systems

"Replicas agree"

In formal logic

"Free from contradiction"

These are the same thing

# 1: Allowing some inconsistencies

Some user-observable inconsistencies might be ok

Two elevators arrive vs. no elevator arrives

Light turns on but no elevator arrives vs. vice versa

Hard to reason about

Need to consider all possible loss/dup/reorder permutations

Need to include node loss too...

*In case of combinatorial explosion, look directly at explosion*

# 2: Collapse loss/dup/reorder space

Want: methodical approach to reduce no. of permutations

Don't care about the order/internals, only the outcomes

Handling:

reordering: receiver is **commutative**

loss: sender retries often enough

duplicates: receiver is **idempotent**

Collapsing all permutations is done by:

Only permuting monotonically ("upward counting")

Removing "resets"/"reverse"/negation generally

Removing this has to be done in the state-feedback loop

This is self-referential negation

**Enforcing consistency is removing the ability to utter paradoxes**

## 2: Collapsing (cont.)

Can solve many (most?) problems:

- Upward-only counters / latest update

- Set/clear order no.  $N$ , then  $N+1$

- Systems that can work with "eventually consistent" data

Sometimes too restrictive



# 3: The middle road

So far, two extremes:

Consider all permutations, or collapse into one

Alternative:

Consider *a few* - how?

Usually only need very targeted reset/reverse

"Reset to initial state", "Go back to 'idle'"

Sometimes caused by external factors

A node that crashes is "forcibly reset"

# Conclusion

Commutativity and idempotency  
are the tools for collapsing the permutation space

Complete permutation collapse is only possible  
if the nodes cannot fail

The amount of protocol you need  
is specification/application dependent



# Distributed systems

