

ENGINEERING

**IT'S LIKE MATH
BUT LOUDER**

Have you tried turning
it off and on again

Defining success and failure

The reasons things go wrong

The expected reasons: External dependencies

- Network (can't connect)

- Files (don't exist)

- Memory (not enough)

- Handled with: exceptions, error returns, etc

The unexpected reasons: Bugs

- "Algorithm" bugs (implementation was incorrect)

- "Specification" bugs (implementation didn't exist)

- Handled with...?

 - Testing? Too late for that now!

The three observers of our error

The developer

- Wants to fix bugs

- Wants fallback strategies when external dependencies fail

The user

- Wants an error message

The program

- Wants to continue

- Wants to ignore errors or self-heal

Our mechanisms should satisfy **all** of these... ideally

When things go wrong

We want unification:

- Detecting all things that go wrong

- Handling/recovering from all things that went wrong

We want little code:

- "Happy path" code should not be swamped by error code

- Don't burden the programmer too much

 - More likely to be written in the first place

 - More likely to be correct

 - Less mental overhead when programming the happy stuff

Identifying the unacceptable

An operation has two outcomes

- It works properly

- It doesn't work properly

Our detection mechanism has two outcomes

- We accept the operation

- We reject the operation

Let's make the obvious figure...

Identifying the unacceptable - 2

		Detection mechanism outcomes	
		We accept it:	We reject it:
Operation outcomes	It works:	Correctly identified successful operations	Failed to identify successful operations
	It doesn't work:	Failed to identify failed operations	Correctly identified failed operations

Identifying the unacceptable - 3

"Traditional" error detection aims to detect failure

Detecting (and rejecting) expected failure is easy

But detecting unexpected failure (bugs) is impossible

Acceptance tests aim to detect success

Detecting (and accepting) expected success is easy

But failing to detect unexpected success is **not** a problem

The path of failure

Error returns

An example of "traditional" error detection

From linux source (cifs_smb3_do_mount)

<https://github.com/torvalds/linux/blob/master/fs/cifs/cifsfs.c#L712-L810>

Notice the appropriate use of goto


Error is returned as data

Control flow and data flow are aligned

Returns vs Exceptions


Always align data flow and control flow

Error returns:



Control flow should go forwards

Exceptions:



Control flow should go **backwards!**

Exceptions (the proper way)

Exceptions are "hard" because try-catch flow goes forward

Use scope/defer instead

```
try {  
    doA();  
    try {  
        doB();  
    } catch(Exception e){  
        undoB();  
        throw e;  
    }  
} catch(Exception e){  
    undoA();  
}
```

```
scope(failure) undoA();  
doA();  
  
scope(failure) undoB();  
doB();
```

Exceptions (the wrong way)

From "Cleaner, more elegant, and wrong"

<https://blogs.msdn.microsoft.com/oldnewthing/20040422-00/?p=39683>

Spot all the things that are wrong!

- Uses try-catch

 - Wrong program flow for the data flow - not the fault of exceptions

- Applies state transformation in incremental steps

 - Multiple sources of errors but only one handler - again, not the fault of exceptions

 - Doesn't split into core/shell

 - Creates some untestable unfathomable Object Oriented abortion

- Duplicates data

 - Either let the inconsistency be ok

 - Or use transactions you fucking bot

 - Again: Not the fault of exceptions

The code paths of failure

The "happy path" uses standard control structures

`if / for / switch / etc`

Errors are when we can't follow the happy path

Use non-standard control structures!

`goto` is not the enemy

Finding success in failure

Writing acceptance tests

Just use `assert()` a lot

(Or `if()` if you want forwards-flow)

Easy to write

Figuring out the logic for them is rarely hard

Low effort

No boilerplate

No extra `err` or `Exception` variables

Doesn't look like dogshit

No extra nesting and curly braces

Acceptance tests example

From the simulator

```
invariant {
    assert(-1 <= currFloor && currFloor < numFloors,
        "currFloor is not between -1..numFloors");
    assert(0 <= prevFloor && prevFloor < numFloors,
        "prevFloor is not between 0..numFloors");
    assert(departDirn != Dirn.Stop,
        "departDirn is Dirn.Stop");
}

(FloorArrival f){
    assert(state.currDirn != Dirn.Stop,
        "Elevator arrived at a floor with currDirn == Dirn.Stop");
    assert(0 <= f && f <= state.numFloors,
        "Elevator \"arrived\" at a non-existent floor\n");
    assert(
        (state.currDirn == Dirn.Up && f >= state.prevFloor) ||
        (state.currDirn == Dirn.Down && f <= state.prevFloor),
        "Elevator arrived at a floor in the opposite direction of travel\n");
    assert(abs(f - state.prevFloor) <= 1,
        "Elevator skipped a floor");
    . . . . .
```

Handling the unacceptable

Acceptance tests that fail just identify that "it didn't work"

The things we get:

- Where it didn't work (file and line, useful for the developer)

- Hopefully an error message written by the developer

The things that are still missing:

- Exactly what originally caused the problem

- A way to recover and heal

How do we self-heal from an arbitrary unknown failure?

Corollary: what's the worst unknown failure?

- Answer: The operation crashes

- (Answer 2: The crash propagates and crashes other things too)

Solution: Just restart it!

Writing restarting-mechanisms

Restarting execution

(Requires "I am alive" messages)

Supervisor

One program is responsible for keeping the other alive

Process pair

The program creates its own duplicate backup

Ensuring data integrity

(Requires that the data passes the acceptance test)

Checkpoint-restart

Copy of data written to persistent storage

Checkpoint-message

"I am alive" message contains data

Persistent

Process operates on a "database" (both read and write)

Integrating with core/shell

The only things that can fail in a process are

- Writing to the outside world

- Your own acceptance tests

We want to preserve the integrity of the process and its outputs

- "Robustness against restarting"

Some guidelines:

- Put the acceptance test in (or right after) the core functions

- Do **all** acceptance tests before **any** state transformation or output

- Stay together

 - Don't apply outputs that must go together in two separate places

 - You are about to fail at doing transactions

 - Don't spread data structures that must go together

Propagation in trees

Overkill

Crashing the entire program is probably overkill

Want to crash only a small part:

- Want other operations to be minimally affected

 - Ongoing operations complete as normal

 - Interactions with restarting operations are delayed but not killed

- Want restarting to be fast

 - Insert the obvious interpretation here

 - Want startup to be un-sequenced so the programmer isn't tempted to create dependency hell

Interactions

If components interact in the "happy path"

Then components must interact during error handling

Minimize error interactions by minimizing **all** interactions

Do **not** invoke behavior in other components

Hey look, it's communicating processes again...

Supervisors

Processes are spawned hierarchically

This gives us a few restarting-strategies:

One-for-one:

Restart just the child process that dies

One-for-all:

If a child dies, restart all children

Rest-for-one:

If a child dies, restart it and the younger ones

Process death needs to be communicated ("I am alive" or similar)

You really want language/library support for this...